

Databases: The other stuff

Transactions, APIs and the other kinds of database

Jo Hallett

December 1, 2025



The story so far...

Two weeks ago we discussed database design

- ▶ We talked about normal forms and doodling designs

Last week we discussed SQL

- ▶ We talked about how to query a database

This week...

- ▶ Well when was the last time you ever heard of someone programming something in SQL only?
- ▶ (...and a bit of logic programming)

Domain specific languages

SQL is a domain specific language

- ▶ It's great for one thing (database queries)
- ▶ But rubbish for general programming
 - ▶ A bit like regular expressions!

You don't want to have to program in SQL all the time

- ▶ So how do you get at it from a more normal programming language?

API bindings

Different programming languages do different things!
SQLite provides bindings for C that are roughly portable

- ▶ Some programming languages will take the C API and make an interface that's more or less the same as the C one
- ▶ Some programming languages will create a brand new API that better fits with the language
- ▶ Some will do something else...

For example Eitaro Fukamachi's SXQL for Commonlisp...

```
(select (:id :name :sex)
  (from (:as :person :p))
  (where (:and (:>= :age 18)
               (:< :age 65)))
  (order-by (:desc :age)))
```

Python and SQLite

Python provides SQLite bindings through the `sqlite3` library.

- ▶ And it's provided as part of a core Python install

```
import sqlite3
con = sqlite3.connect("chinook.db")
cur = con.cursor()
for row in cur.execute("SELECT * FROM Album LIMIT 3"):
    print(row)
cur.close()
con.close()
```

```
(1, 'For Those About To Rock We Salute You', 1)
(2, 'Balls to the Wall', 2)
(3, 'Restless and Wild', 2)
```

Well except...

What if something goes wrong in the middle?

- ▶ I might crash with my database still open

```
import sqlite3
try:
    con = sqlite3.connect("chinook.db")
    try:
        cur = con.cursor()
        for row in cur.execute("SELECT * FROM Album LIMIT 3"):
            print(row)
    finally:
        cur.close()
finally:
    con.close()
```

```
(1, 'For Those About To Rock We Salute You', 1)
(2, 'Balls to the Wall', 2)
(3, 'Restless and Wild', 2)
```

What about if I'm adding data?

What happens if I'm adding multiple bits of data

```
import sqlite3
try:
    con = sqlite3.connect("chinook.db")
    try:
        artist = (276, "Belle_and_Sebastian")
        albums = [(348, "Dear_Catastrophe_Waitress", 276),
                   (349, "Write_About_Love", 276)]
        cur = con.cursor()
        cur.execute("INSERT INTO Artist(ArtistId, Name) VALUES(?,?)", artist)
        for album in albums:
            cur.execute('''
                INSERT INTO Album(AlbumId, Title, ArtistId)
                VALUES(?, ?, ?)'',
                        album)
        #con.commit() #Uncomment if you want to save changes!
        for row in cur.execute(''
            SELECT Artist.Name, Album.Title
            FROM Album JOIN Artist ON Album.ArtistId = Artist.ArtistId
            WHERE Album.ArtistId = 276''):
            print(row)
    finally:
        cur.close()
finally:
    con.close()
```

None

??

What are all those ?'s about?

```
cur.execute("INSERT\u2014INTO\u2014Artist(ArtistId,\u2014Name)\u2014VALUES\u2014(?,?)", artist)
```

Why not?

```
cur.execute("INSERT\u2014INTO\u2014Artist(ArtistId,\u2014Name)\u2014 +  
    "VALUES\u2014("+artist[0]+",\"artist[1]\")")
```

Little Johnny Drop Tables

Suppose you have a login system on a website.

```
cur.execute("SELECT FROM users WHERE username = '" + username + "' AND password = '" + password + "'")
```

If your database can find a row that matches their username and password they get to log in!
My password is fish' OR 'fish' = 'fish

- ▶ Good secure password
- ▶ What will happen if I try and login?

```
cur.execute("SELECT FROM users WHERE username = 'jo' OR password = 'fish'")
```

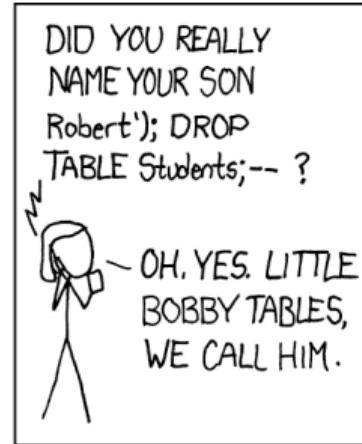
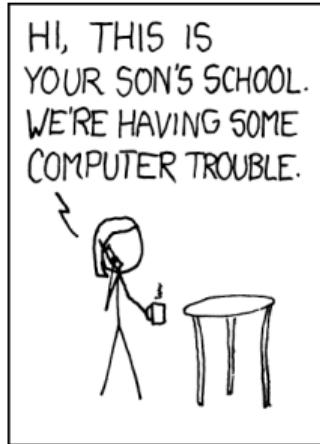
What would happen if Matt tried to log in with my password?

Prepared Statements

SQL Injection vulnerabilities

Prepared statements are a safety measure

- ▶ When I insert data it goes through what I've added and escapes it



Don't get your database hacked

- ▶ Use prepared statements
- ▶ ...and don't store passwords like this

I'll commit to that

```
#con.commit() #Uncomment if you want to save changes!
```

Are my changes not normally saved?

What if something goes wrong?

Suppose I'm adding lots of albums and artists...

- ▶ I get halfway through and then my program crashes
- ▶ What happens if I try and restart?

```
.schema Artist
```

```
CREATE TABLE [Artist]
(
    [ArtistId] INTEGER NOT NULL,
    [Name] NVARCHAR(120),
    CONSTRAINT [PK_Artist] PRIMARY KEY ([ArtistId])
);
```

A UNIQUE sort of pain

If I restart I have to:

- ▶ Work out how far I got and undo all bits I've already done
- ▶ Work out how far I got and not skip some of the entries to do

Both sound like a pain.

Transactions

- ▶ Let me define a bunch of database queries as a single atomic operation
- ▶ If any fail I can roll back all the changes and start again

Until I `commit()` my work nothing is saved

Well, technically you can do this from SQL directly...

```
BEGIN TRANSACTION;

INSERT INTO Artist(ArtistId, Name)
VALUES (4949, "Geordie\u00e2Greep");
INSERT INTO Album(AlbumId, Title, ArtistId)
VALUES (12345, "The\u00e2New\u00e2Sound", 4949);

ROLLBACK TRANSACTION; -- or...
COMMIT;
```

But unless you work with SQL directly, it's not that common...

Java and the JDBC

With Python, we loaded SQLite's API

- ▶ Java tries something different

Rather than implementing SQL API's for every database engine...

- ▶ Implement a general framework for database access
 - ▶ The Java Database Connectivity or JDBC
- ▶ Database engines provide their own bridge to the JDBC API

So how do I use it?

```
import java.sql.*;

try {
    final var connection = DriverManager.getConnection("jdbc:sqlite:chinook.db");
    final var statement = connection.prepareStatement("SELECT name FROM Artist WHERE name LIKE ?");
} {
    statement.setString(1, "%B%");
    final var results = statement.executeQuery();
    while (results.next()) {
        System.out.println(results.getObject(0));
    }
} catch (SQLException err) {
    err.printStackTrace(System.err);
}
```

java.sql.SQLException: No suitable driver found for jdbc:sqlite:chinook.db

And with the CLASSPATH fixed...

Go and fetch the driver you need (or use Maven) and stick it in your classpath

```
import java.sql.*;

try {
    final var connection = DriverManager.getConnection("jdbc:sqlite:chinook.db");
    final var statement = connection.prepareStatement("SELECT name FROM Artist WHERE name LIKE ?");
} {
    statement.setString(1, "Q%");
    final var results = statement.executeQuery();
    while (results.next()) {
        System.out.println(results.getObject(1));
    }
} catch (SQLException err) {
    err.printStackTrace(System.err);
}
```

Queen

That's all folks

We've covered:

- ▶ Some database theory
- ▶ SQL syntax
- ▶ Interacting with Databases from programming languages

Well except...

Lets make a tree in SQL!

```
CREATE TABLE tree
( id INTEGER PRIMARY KEY AUTOINCREMENT
, value TEXT NOT NULL
, left INTEGER
, right INTEGER
);
```

Assuming your table contains multiple trees:

- ▶ Can you write a query to search if a given value is in a tree or not?

Recursive data structures

Turns out you can't.

- ▶ Well you can, but you have to write out queries for each level of nesting.

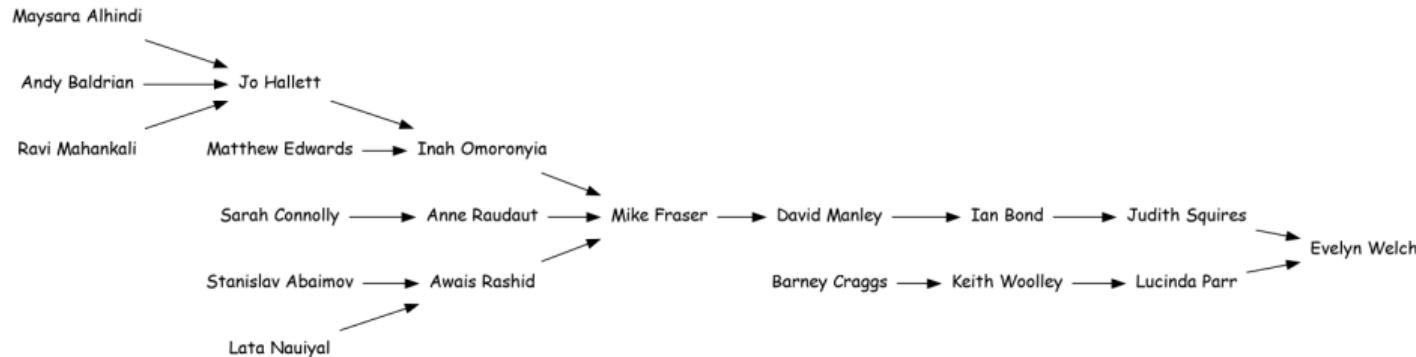
```
SELECT *
FROM tree
JOIN tree AS tree2
ON tree.left = tree2.id
JOIN tree AS tree3
ON tree2.left = tree3.id
JOIN tree AS tree4
ON tree3.left = tree4.id
JOIN tree AS tree5
ON tree4.left = tree5.id
JOIN tree AS tree6
ON tree5.left = tree6.id
JOIN tree AS tree7
ON tree6.left = tree7.id
-- and so on...
```

This is tedious.

How do we express these in SQL

Trees and other graph structures turn up a lot in real life; but SQL can't handle them in general
For example:

- ▶ Suppose you are storing a database of employees supervisors...
- ▶ To have your holiday approved you need approval from your boss, or their boss or their boss's boss...



I told Barney and Awais that I'm going on holiday today; and they said it was fine...

- ▶ Am I allowed to take the day off?

I could do this in SQL...

```
import sqlite3
person = "Jo\u00e3oHallett"
boss1 = "Awais\u00e3oRashid"
boss2 = "Barney\u00e3oCraggs"

def canAuthorizeLeave(person, boss):
    con = sqlite3.connect("bosses.db")
    cur = con.cursor()
    for row in cur.execute("SELECT\u00a0manager\u00a0FROM\u00a0manager\u00a0WHERE\u00a0employee\u00a0IS\u00a0?", [person]):
        if row is None:
            return False
        elif row[0] == boss:
            return True
        else:
            return canAuthorizeLeave(row[0], boss)
    cur.close()
    con.close()

if canAuthorizeLeave(person, boss1) or canAuthorizeLeave(person, boss2):
    print("Leave\u00a0approved!")
else:
    print("Go\u00a0ask\u00a0someone\u00a0else!")
```

Go ask someone else!

Yuk

But it feels clunky...

- ▶ Can't do from SQL alone
- ▶ Have to make repeated queries
- ▶ Surprisingly large amount of boilerplate for one line of database query

The thing is if we were to write this out in formal logic its relatively trivial...

$$\text{R1} \quad \frac{\Gamma \vdash \text{manages}(\text{Person}, \text{Boss})}{\Gamma \vdash \text{canAuthorizeLeave}(\text{Person}, \text{Boss})}$$

$$\text{R2} \quad \frac{\Gamma \vdash \text{manages}(\text{Person}, \text{Boss}) \quad \Gamma \vdash \text{manages}(\text{Boss}, \text{BossesBoss})}{\Gamma \vdash \text{manages}(\text{Person}, \text{BossesBoss})}$$

Proof!

I could even generate a proof tree to prove it...

What if I had asked Ian...

- ▶ Read $\Gamma \vdash$ as we know that...
- ▶ (Or more accurately: Γ contains statements from which we can derive that...)

$$\frac{\Gamma \vdash \text{manages(jo, inah)}}{\Gamma \vdash \text{canAuthorizeLeave(jo, ian)}}$$
$$\frac{\Gamma \vdash \text{manages(inah, mike)}}{\frac{\Gamma \vdash \text{manages(mike, david)} \quad \Gamma \vdash \text{manages(david, ian)}}{\frac{\Gamma \vdash \text{manages(mike, ian)}}{\Gamma \vdash \text{manages(inah, ian)}}}}$$

(I was gonna ask Evelyn to authorize my leave but the tree got too big to fit on the slide :-S)

If only we had a database language based on first order logic...

Datalog

Datalog is a database language based on first order logic

- ▶ Simplified version of the logic programming language Prolog
- ▶ All facts/assertions are written as Horn clauses
 - ▶ predicate(arguments)
- ▶ Variables are capitalized, constants are not
- ▶ Rules separated with a `:``-`, conjunctions with `,`
 - ▶ `canAuthorizeLeave(P,B) :- manages(P,B).`
 - ▶ `manages(P,BB) :- manages(P,B), manages(B,BB).`
- ▶ All variables in the left hand of a rule must be referred to in the right hand side.
 - ▶ Typing is usually enough to always meet this rule...
 - ▶ `canRead(U,F) :- user(U), file(F), otherReadBitSet(F).`
 - ▶ Cannot express infinite sized sets (e.g. integers)
- ▶ Closed world assumption
 - ▶ The database contains all the facts and rules you might ever need to prove something
 - ▶ If you can't prove it with all the facts then it's false

Pros and Cons

Datalog is good because:

- ▶ Can represent trees
- ▶ Can provably find all possibilities in a reasonable time
- ▶ Can produce proofs given the right implementation
- ▶ With the right implementation can find all values of a variable satisfying a query
- ▶ With extensions can run quite complex queries

Datalog is bad a pain because:

- ▶ No infinite sets (no numbers)
- ▶ Research-grade implementations
- ▶ Somewhat old fashioned
 - ▶ A tool from the dark ages of AI...

But its making a bit of a comeback?

- ▶ AI is cool again?

Applications

Most databases are written in SQL

- ▶ But sometimes Datalog is a really powerful tool

Applications include:

Program analysis can you prove this property about a program?

Access control can this person do this action

Machine learning and AI automated reasoning

Google/Facebook-scale data you do not have yottabytes of data...

Access control and Datalog

I love access control papers and spent my PhD working on them

- ▶ (Well a variant of an authorization logic called SecPAL)

If you read an access control paper you will see inference rules

- ▶ Turns out knowing that you'll be able to make a decision in a reasonable period of time and with a proof of why it is correct is really good for access control...

Let's implement the UNIX DAC!

- ▶ You can read a file if you own it and the user read permission is set
- ▶ You can read a file if its owned by a group you're in and the group read permission is set
- ▶ You can read a file if the other read permission is set

First order logic

You can read a file if you own it and the user read permission is set

User

$$\frac{\Gamma \vdash U \text{ owns } F \quad \Gamma \vdash F \text{ has user read bit set}}{\Gamma \vdash U \text{ can read } F}$$

You can read a file if its owned by a group you're in and the group read permission is set

Group

$$\frac{\Gamma \vdash G \text{ owns } F \quad \Gamma \vdash U \text{ member of } G \quad \Gamma \vdash F \text{ has group read bit set}}{\Gamma \vdash X \text{ can read } F}$$

You can read a file if the other read permission is set

Other

$$\frac{\Gamma \vdash F \text{ has other read bit set}}{\Gamma \vdash X \text{ can read } F}$$

Given...

User

$$\frac{\Gamma \vdash U \text{ owns } F \quad \Gamma \vdash F \text{ has user read bit set}}{\Gamma \vdash U \text{ can read } F}$$

Group

$$\frac{\Gamma \vdash G \text{ owns } F \quad \Gamma \vdash U \text{ member of } G \quad \Gamma \vdash F \text{ has group read bit set}}{\Gamma \vdash U \text{ can read } F}$$

Other

$$\frac{\Gamma \vdash F \text{ has other read bit set}}{\Gamma \vdash U \text{ can read } F}$$

$\Gamma \vdash \text{Matt owns .ssh/id_ed25519}$

$\Gamma \vdash \text{.ssh/id_ed25519 has other read bit set}$

Can I read Matt's private key?

And in action?

```
file('.ssh/id_ed25519').  
file('jos-diary').  
user('matt'). user('jo'). group('users').  
  
canRead(U,F):- user(U), file(F),  
    owns(U,F), userReadBitSet(F).  
  
canRead(U,F):- user(U), group(G), file(F),  
    member(U,G), owns(G,F), groupReadBitSet(F).  
  
canRead(U,F):- user(U), file(F),  
    otherReadBitSet(F).  
  
owns('matt','.ssh/id_ed25519').  
owns('jo','jos-diary').  
member('matt','users').  
member('jo','users').  
userReadBitSet('.ssh/id_ed25519').  
userReadBitSet('jos-diary').  
otherReadBitSet('.ssh/id_ed25519').
```

I said in action?

```
$ swipl output.pl
?- trace.
true.

[trace] ?- canRead('jo', '.ssh/id_ed25519').
Call: (12) canRead(jo, '.ssh/id_ed25519') ? creep
Call: (13) user(jo) ? creep
Exit: (13) user(jo) ? creep
Call: (13) file('.ssh/id_ed25519') ? creep
Exit: (13) file('.ssh/id_ed25519') ? creep
Call: (13) owns(jo, '.ssh/id_ed25519') ? creep
Fail: (13) owns(jo, '.ssh/id_ed25519') ? creep
Redo: (12) canRead(jo, '.ssh/id_ed25519') ? creep
Call: (13) user(jo) ? creep
Exit: (13) user(jo) ? creep
Call: (13) group(_21294) ? creep
Exit: (13) group(users) ? creep
Call: (13) file('.ssh/id_ed25519') ? creep
Exit: (13) file('.ssh/id_ed25519') ? creep
Call: (13) member(jo, users) ? creep
Exit: (13) member(jo, users) ? creep
Call: (13) owns(users, '.ssh/id_ed25519') ? creep
Fail: (13) owns(users, '.ssh/id_ed25519') ? creep
Redo: (12) canRead(jo, '.ssh/id_ed25519') ? creep
Call: (13) user(jo) ? creep
Exit: (13) user(jo) ? creep
Call: (13) file('.ssh/id_ed25519') ? creep
Exit: (13) file('.ssh/id_ed25519') ? creep
Call: (13) otherReadBitSet('.ssh/id_ed25519') ? creep
Exit: (13) otherReadBitSet('.ssh/id_ed25519') ? creep
Exit: (12) canRead(jo, '.ssh/id_ed25519') ? creep
true.
```

Are you ever really going to use Datalog?

Probably not

- ▶ But it's a cool tool
- ▶ And what's the point of teaching you Software Tools if I can't occasionally show you the weird ones?

But maybe...

- ▶ Maybe one day you'll be building a proof tool
- ▶ Or a recursive database
- ▶ And you'll be finding SQL a pain
- ▶ And you'll remember this weird little tool...

Recap

- ▶ You can talk to SQL from normal programming languages
- ▶ Math sometimes gets implemented as a programming language...
- ▶ Man I hate computers...